

# Lecture 6: Inference Pat 1

谢丹  
清华大学数学系

November 9, 2025

# Clustering Methods: Overview and Preview I

## Previously Covered Methods

We have explored two fundamental clustering approaches:

### K-Means Clustering

- ▶ Starts with  $K$  initial centroids
- ▶ **Iteratively updates:**
  - ▶ Cluster assignments
  - ▶ Centroid positions
- ▶ Based on **minimum distance** criterion

### Gaussian Mixture Models (GMM)

- ▶ Probabilistic model with parameters:
  - ▶ Mixing coefficients:  $\pi_k$
  - ▶ Means:  $\mu_k$
  - ▶ Covariances:  $\Sigma_k$
- ▶ **Iteratively updates** parameters via EM algorithm
- ▶ Based on **probability density** estimation

# Clustering Methods: Overview and Preview II

## Common Assumption

Both methods assume clusters are centered around specific points (centroid-based).

## Next: Alternative Clustering Paradigms

We will now explore two fundamentally different clustering approaches that overcome limitations of centroid-based methods.

# Hierarchical Clustering: Overview

## Core Idea

Build a tree-like hierarchy of clusters where clusters at higher levels contain clusters from lower levels.

## Two Main Approaches

### 1. Agglomerative (Bottom-up)

- ▶ Start with each point as a cluster
- ▶ Iteratively merge closest clusters

### 2. Divisive (Top-down, less common)

- ▶ Start with one cluster
- ▶ Recursively split clusters

# Agglomerative Hierarchical Clustering Algorithm

## Step-by-Step Process

1. Start with  $n$  clusters (each point is its own cluster)
2. Compute  $n \times n$  proximity matrix
3. Repeat until one cluster remains:
  - ▶ Find two closest clusters  $C_i$  and  $C_j$
  - ▶ Merge  $C_i$  and  $C_j$
  - ▶ Update proximity matrix

## Proximity Matrix Update

After merging clusters  $C_i$  and  $C_j$  into  $C_{ij}$ , update distances to other clusters  $C_k$  using linkage method.

# Linkage Methods

## Single Linkage

Minimum distance between clusters:

$$d(C_i, C_j) = \min_{a \in C_i, b \in C_j} d(a, b)$$

## Complete Linkage

Maximum distance between clusters:

$$d(C_i, C_j) = \max_{a \in C_i, b \in C_j} d(a, b)$$

## Average Linkage

Average distance between clusters:

$$d(C_i, C_j) = \frac{1}{|C_i||C_j|} \sum_{a \in C_i} \sum_{b \in C_j} d(a, b)$$

## Ward's Method

Minimize increase in within-cluster variance:


$$d(C_i, C_j) = \frac{|C_i||C_j|}{|C_i| + |C_j|} \|\mu_i - \mu_j\|^2$$

	$P_1$	$P_2$	$P_3$	$P_4$	$P_5$	$P_6$
$P_1$	0					
$P_2$	1	0				
$P_3$	2.5	2.8	0			
$P_4$	3	3.5	1.2	0		
$P_5$	2	2.3	1.5	1.8	0	
$P_6$	5	6	2.6	2.1	3.2	0


	$C_1$	$C_2$	$C_3$	$C_4$	$C_5$	$C_6$
$C_1$	0					
$C_2$	2.2	0				
$C_3$	2.8	1.2	0			
$C_4$	1.8	1.5	1.7	0		
$C_5$	6	2.6	2.1	3.2	0	
$C_6$						




	$c_1$	$c_2$	$p_5$	$p_4$
$c_1$	0			
$c_2$	2.7	0		
$p_5$	1.8	1.4	0	
$p_4$	6	2.1	3.2	0



	$c_1$	$c_3$	$p_6$
$c_1$	0		
$c_3$	2	0	
$p_6$	6	1.8	0



	$C_L$	$C_R$
$C_L$	C	
$C_R$		D



The optimal number of clusters can be found using a dendrogram (the last diagram in last slide): Find the longest vertical line and cut it.

# DBSCAN: Core Concepts

## Philosophy

Cluster based on density connectivity rather than distance from centroids.  $MinPts$  is the number of minimal points.

## Key Definitions

- ▶  **$\varepsilon$  -neighborhood:**  $N_\varepsilon(p) = \{q \in D | dist(p, q) \leq \varepsilon\}$
- ▶ **Core point:**  
 $|N_\varepsilon(p)| \geq MinPts$
- ▶ **Border point:** In  $\varepsilon$  of core point but  
 $|N_\varepsilon(p)| < MinPts$
- ▶ **Noise point:** Neither core nor border

# DBSCAN: Reachability Concepts

## Directly Density-Reachable

Point  $q$  is directly density-reachable from  $p$  if:

- ▶  $p$  is a core point
- ▶  $q \in N_\epsilon(p)$

## Density-Reachable

Point  $q$  is density-reachable from  $p$  if there exists a chain  $p_1, \dots, p_n$  with  $p_1 = p$ ,  $p_n = q$  where each  $p_{i+1}$  is directly density-reachable from  $p_i$ .

## Density-Connected

Points  $p$  and  $q$  are density-connected if there exists point  $o$  such that both  $p$  and  $q$  are density-reachable from  $o$ .




Figure: Density connectivity

# Simplified DBSCAN Algorithm I

## Step 1 — Identify Point Types

Identify all points as either:

- ▶ **Core point**: Has at least MinPts within  $\varepsilon$  distance
- ▶ **Border point**: Within  $\varepsilon$  of a core point but has  $<$  MinPts neighbors
- ▶ **Noise point**: Neither core nor border point

## Step 2 — Process Core Points

For each unclustered core point:

**2a** Create a new cluster

**2b** Expand cluster by adding all points that are:

- ▶ Unclustered
- ▶ Density-connected to the current core point

## Step 3 — Assign Border Points

# Simplified DBSCAN Algorithm II

For each unclustered border point, assign it to the cluster of the nearest core point.

## Step 4 — Handle Noise Points

Leave all noise points unclustered (marked as outliers).

## Final Result

- ▶ Dense regions form clusters of arbitrary shapes
- ▶ Sparse regions are identified as noise
- ▶ No need to pre-specify number of clusters

# DBSCAN vs Hierarchical Clustering: Comparison


Hierarchical Clustering	DBSCAN
<b>Cluster Shape</b> Tends to find spherical clusters Shape determined by linkage	<b>Cluster Shape</b> Finds arbitrarily shaped clusters Shape follows density contours
<b>Noise Handling</b> Poor - sensitive to outliers Outliers can distort entire structure	<b>Noise Handling</b> Excellent - explicit noise category Noise points don't affect clusters
<b>Parameters</b> Linkage method, distance metric Number of clusters (for cutting)	<b>Parameters</b> $\varepsilon$ (eps), MinPts No need for cluster count
<b>Complexity</b> $O(n^3)$ naive, $O(n^2 \log n)$ optimized $O(n^2)$ space	<b>Complexity</b> $O(n^2)$ naive, $O(n \log n)$ with indexing $O(n)$ space
<b>Output</b> Dendrogram (hierarchical structure)	<b>Output</b> Flat partition + noise points

# CHAPTER 7:

# Inference

# Explain Away Phenomenon: Three-Node Example

Let's first explain an interesting fact about the inference.



## Bayesian Network Structure

- ▶ Two independent causes that can produce the same effect
- ▶ This is a **v-structure** or **collider**:  
 $Rain \rightarrow WetGrass \leftarrow Sprinkler$
- ▶ The causes are independent *a priori* but become dependent when we observe the effect

# Probability Model Specification

## Prior Probabilities

- ▶  $P(\text{Rain} = T) = 0.2$
- ▶  $P(\text{Sprinkler} = T) = 0.1$
- ▶  $\text{Rain} \perp \text{Sprinkler}$

## Key Insight

- ▶ Initially:  $\text{Rain} \perp \text{Sprinkler}$
- ▶ After observing  $\text{WetGrass} = T$ :  
 $\text{Rain} \not\perp \text{Sprinkler} | \text{WetGrass}$
- ▶ The causes become **negatively correlated**

## Conditional Probability Table

$P(\text{WetGrass} | \text{Rain}, \text{Sprinkler})$ :

Rain	Sprinkler	WetGrass	Probability
F	F	T	0.0
F	F	F	1.0
F	T	T	0.9
F	T	F	0.1
T	F	T	0.8
T	F	F	0.2
T	T	T	0.98
T	T	F	0.02

# The "Explain Away" Effect Step by Step

## Step 1: Initial Beliefs (No Evidence)

- ▶  $P(Rain = T) = 0.2$
- ▶  $P(Sprinkler = T) = 0.1$
- ▶  $P(Rain = T \cap Sprinkler = T) = 0.2 \times 0.1 = 0.02$

## Step 2: Observe Wet Grass

We observe:  $WetGrass = True$

What happens to our beliefs about the causes?

## Step 3: Updated Beliefs

Using Bayes' theorem:

$$P(R|W) = \frac{P(W|R)P(R)}{P(W)}$$

$$P(S|W) = \frac{P(W|S)P(S)}{P(W)}$$

# Calculating the Evidence

## Probability of Wet Grass

$$\begin{aligned} P(W = T) &= \sum_{r,s} P(W = T|r, s)P(r)P(s) \\ &= P(W = T|R = F, S = F)P(R = F)P(S = F) \\ &\quad + P(W = T|R = F, S = T)P(R = F)P(S = T) \\ &\quad + P(W = T|R = T, S = F)P(R = T)P(S = F) \\ &\quad + P(W = T|R = T, S = T)P(R = T)P(S = T) \\ &= 0.0 \times 0.8 \times 0.9 + 0.9 \times 0.8 \times 0.1 \\ &\quad + 0.8 \times 0.2 \times 0.9 + 0.98 \times 0.2 \times 0.1 \\ &= 0 + 0.072 + 0.144 + 0.0196 = 0.2356 \end{aligned}$$

# Posterior Probabilities

## Individual Posteriors

$$\begin{aligned}P(R|W) &= \frac{P(W|R)P(R)}{P(W)} \\&= \frac{0.8 \times 0.2}{0.2356} \approx 0.679\\P(S|W) &= \frac{P(W|S)P(S)}{P(W)} \\&= \frac{0.9 \times 0.1}{0.2356} \approx 0.382\end{aligned}$$

### Key Observation

- ▶ Both probabilities increased:  
 $P(R)$ :  $0.2 \rightarrow 0.679$   
 $P(S)$ :  $0.1 \rightarrow 0.382$
- ▶ But wait... there's more!

# The "Explain Away" Revealed

## Joint Posterior Probability

What is  $P(R = T, S = T | W = T)$ ?

$$\begin{aligned} P(R = T, S = T | W = T) &= \frac{P(W = T | R = T, S = T)P(R = T)P(S = T)}{P(W = T)} \\ &= \frac{0.98 \times 0.2 \times 0.1}{0.2356} \\ &= \frac{0.0196}{0.2356} \approx 0.083 \end{aligned}$$

## The Explain Away Effect!

- ▶ If independent:

$$P(R, S | W) = P(R | W)P(S | W) \approx 0.679 \times 0.382 \approx 0.259$$

- ▶ Actual:  $P(R, S | W) \approx 0.083$
- ▶ **They are much less likely to co-occur than expected!**

# Mathematical Generalization

## General Common Effect Structure

For any v-structure:  $A \rightarrow C \leftarrow B$

- ▶ Initially:  $A \perp B$
- ▶ After observing  $C$ :  $A \not\perp B|C$
- ▶ The correlation is generally negative

## Bayes' Theorem Explanation

$$\begin{aligned} P(A, B|C) &= \frac{P(C|A, B)P(A)P(B)}{P(C)} \\ &\propto P(C|A, B)P(A)P(B) \end{aligned}$$

The conditional probability  $P(C|A, B)$  creates the dependence.

## Key Insight for Inference

- ▶ This is why inference algorithms must carefully handle v-structures

# Summary

## Key Points

- ▶ **Explain away** occurs in common effect structures (v-structures)
- ▶ Independent causes become dependent when their common effect is observed
- ▶ The dependence is typically **negative correlation**
- ▶ One cause can "explain away" the need for the other cause


## Importance in Probabilistic Inference

- ▶ Crucial for correct belief updating in Bayesian networks
- ▶ Explains counter-intuitive reasoning patterns
- ▶ Fundamental to understanding conditional independence
- ▶ Essential for efficient inference algorithms

## **Section 1: Exact inference**

# The Chain Graph Structure

Bayesian Network



Joint Probability Factorization

$$P(\mathbf{X}) = P(X_1) \prod_{i=2}^N P(X_i | X_{i-1})$$

Inference Goal

Compute marginals:

$$P(X_i) = \sum_{\mathbf{X} \setminus X_i} P(\mathbf{X})$$

# Key Insight: Exploiting Factorization I

## Factorization Enables Efficient Computation

We observe that the joint probability distribution admits a **factorized form**, allowing us to **exchange the order of summation** when computing marginals. This crucial insight leads to a dramatic reduction in computational complexity.

### Example

Sequential Marginalization Strategy Instead of summing over all variables simultaneously, we can proceed **sequentially**:

- ▶ First marginalize the last node, effectively removing it from consideration
- ▶ Then proceed to the next node, leveraging the simplified structure
- ▶ Continue this process until reaching the target variable

## General Framework: Factor Graphs

# Key Insight: Exploiting Factorization II

For arbitrary graphical models, we formalize this approach using **factor graphs**:


- ▶ Explicitly represent the factorization structure
- ▶ Enable systematic message-passing protocols
- ▶ Provide a unified framework for efficient inference

## Complexity Reduction

This strategy transforms the computation from **exponential** to **polynomial** complexity, making exact inference feasible for large-scale models!

# Factor Graph Representation I

Factor Graph consists of variables and the factors in the joint probability



$$f_1(X_1) = P(X_1)$$

$$f_2(X_1, X_2) = P(X_2|X_1)$$

$$f_3(X_2, X_3) = P(X_3|X_2)$$

⋮

$$f_N(X_{N-1}, X_N) = P(X_N|X_{N-1})$$

# Factor Graph Representation II

## Key Insight

Factorization enables efficient local computations via message passing

# Message Passing: Forward Pass ( $\alpha$ ) |




Forward Messages

$$\alpha_2(X_2) = \sum_{x_1} f_1(x_1) f_2(x_1, X_2)$$

$$\alpha_3(X_3) = \sum_{x_2} \alpha_2(x_2) f_3(x_2, X_3)$$

$$\alpha_i(X_i) = \sum_{x_{i-1}} \alpha_{i-1}(x_{i-1}) f_i(x_{i-1}, X_i)$$

## Message Passing: Backward Pass ( $\beta$ )



Backward Messages

$$\beta_{N-1}(X_{N-1}) = \sum_{x_N} f_N(X_{N-1}, x_N)$$

$$\beta_{N-2}(X_{N-2}) = \sum_{x_{N-1}} f_{N-1}(X_{N-2}, x_{N-1}) \beta_{N-1}(x_{N-1})$$

$$\beta_i(X_i) = \sum_{x_{i+1}} f_{i+1}(X_i, x_{i+1}) \beta_{i+1}(x_{i+1})$$

# Marginal Computation I

## Combining Messages

For any variable  $X_i$ , combine forward and backward information:

$$P(X_i) \propto \alpha_i(X_i) \cdot \beta_i(X_i)$$

## Normalization

$$P(X_i) = \frac{\alpha_i(X_i)\beta_i(X_i)}{\sum_{x_i} \alpha_i(x_i)\beta_i(x_i)}$$

## Example

# Marginal Computation II

## 3-Node Chain

$$P(X_2) \propto \alpha_2(X_2) \cdot \beta_2(X_2)$$

$$\alpha_2(X_2) = \sum_{x_1} P(x_1)P(X_2|x_1)$$

$$\beta_2(X_2) = \sum_{x_3} P(x_3|X_2)$$

# Complete Two-Pass Algorithm

## Step 1: Forward Pass

1. Initialize:  $\alpha_1(X_1) = P(X_1)$
2. For  $i = 2$  to  $N$ :  $\alpha_i(X_i) = \sum_{x_{i-1}} \alpha_{i-1}(x_{i-1})P(X_i|x_{i-1})$

## Step 2: Backward Pass

1. Initialize:  $\beta_N(X_N) = 1$
2. For  $i = N - 1$  to  $1$ :  $\beta_i(X_i) = \sum_{x_{i+1}} P(x_{i+1}|X_i)\beta_{i+1}(x_{i+1})$

## Step 3: Compute Marginals

For  $i = 1$  to  $N$ :

$$P(X_i) = \frac{\alpha_i(X_i)\beta_i(X_i)}{\sum_{x_i} \alpha_i(x_i)\beta_i(x_i)}$$

# Computational Complexity

## Complexity Analysis

- ▶ **Naive approach:**  $O(K^N)$
- ▶ **Sum-Product:**  $O(NK^2)$
- ▶ **Speedup:** Exponential to polynomial!

## Key Operations per Step

- ▶ Message computation:  
 $O(K^2)$
- ▶  $N - 1$  messages in each direction
- ▶ Total:  $O(NK^2)$  operations

## Memory Requirements

- ▶ Store all  $\alpha_i$ :  $O(NK)$
- ▶ Store all  $\beta_i$ :  $O(NK)$
- ▶ Total:  $O(NK)$  storage

## Optimal for Chains

Sum-Product is provably optimal for exact inference in chain graphs

# The General Inference Problem

- ▶ Joint probability distribution over variables  
 $\mathbf{X} = \{X_1, \dots, X_N\}$ :

$$P(\mathbf{X}) = \frac{1}{Z} \prod_c \psi_c(\mathbf{X}_c)$$

- ▶  $\psi_c$ : Factor/potential function (both undirected and directed graph model)
- ▶  $Z$ : Normalization constant

## Goal

Compute marginal distribution for variable  $X_i$ :

$$P(X_i) = \sum_{\mathbf{X} \setminus X_i} P(\mathbf{X})$$

## Challenge

Naive summation has exponential complexity! Sum-Product algorithm provides efficient solution for **tree-structured graphs**.

# Factor Graph Representation


## Two Types of Nodes

- ▶ **Variable Nodes** (circles)
- ▶ **Factor Nodes** (squares)

## Example

$P =$

$$f_A(X_1)f_B(X_1, X_2)f_C(X_2, X_3)f_D(X_3)$$



## Why Factor Graphs?

- ▶ Clear representation of factorization
- ▶ Explicit message passing rules
- ▶ Generalizes Bayesian Networks and MRFs

# Message Passing Rules I


## Two Message Types

1.  $\mu_{X \rightarrow f}(X)$ : Variable  $\rightarrow$  Factor
2.  $\mu_{f \rightarrow X}(X)$ : Factor  $\rightarrow$  Variable

## Leaf Initialization

- ▶ Variable leaf:  $\mu_{X \rightarrow f}(X) = 1$
- ▶ Factor leaf:  
 $\mu_{f \rightarrow X}(X) = f(X)$

## Message Rules



# Message Passing Rules II

- ▶ **Variable → Factor:**  $\mu_{X \rightarrow f}(X) = \prod_{h \in \text{ne}(X) \setminus \{f\}} \mu_{h \rightarrow X}(X)$
- ▶ **Factor → Variable:**

$$\mu_{f \rightarrow X}(X) = \sum_{\mathbf{Y}} \left( f(\mathbf{Y}, X) \cdot \prod_{Y \in \text{ne}(f) \setminus \{X\}} \mu_{Y \rightarrow f}(Y) \right)$$

# General Algorithm Steps

## Step 1: Choose Root

Pick an arbitrary node as the root of the tree.

## Step 2: Leafward Pass (Collection)

- ▶ Start from leaves, pass messages inward toward root
- ▶ Node sends message after receiving from all children

## Step 3: Rootward Pass (Distribution)

- ▶ Root sends messages back toward leaves
- ▶ Complete message propagation in both directions

## Step 4: Compute Marginals

For each variable  $X_i$ :

$$P(X_i) = \frac{1}{Z} \prod_{f \in \text{ne}(X_i)} \mu_{f \rightarrow X_i}(X_i)$$

# Worked Example

## Factorization

$$P(X_1, X_2, X_3) = f_A(X_1)f_B(X_1, X_2)f_C(X_2, X_3)$$

Compute  $P(X_2)$

$$\mu_{X_1 \rightarrow f_B} = f_A(X_1)$$

$$\mu_{X_3 \rightarrow f_C} = 1$$

$$\mu_{f_C \rightarrow X_2} = \sum_{x_3} f_C(X_2, x_3)$$

$$\mu_{f_B \rightarrow X_2} = \sum_{x_1} f_B(x_1, X_2)f_A(x_1)$$



Final Marginal

$$P(X_2) \propto \mu_{f_B \rightarrow X_2} \cdot \mu_{f_C \rightarrow X_2}$$

$$P(X_2) \propto \left[ \sum_{x_1} f_A f_B \right] \cdot \left[ \sum_{x_3} f_C \right]$$

# Key Properties and Applications

## Computational Efficiency

- ▶ **Naive:**  $O(K^N)$
- ▶ **Sum-Product:** Linear in nodes
- ▶ Cost per message:  $O(K^m)$  for factor with  $m$  variables

## Theoretical Guarantees

- ▶ **Exact** for tree-structured graphs
- ▶ Computes **all marginals** efficiently
- ▶ Foundation for many algorithms

## Applications and Extensions

- ▶ **Forward-Backward Algorithm** (HMMs)
- ▶ **Kalman Filter**
- ▶ **Loopy Belief Propagation** for graphs with cycles
- ▶ **Expectation Propagation**

## Limitation

Exact inference only for **trees**.  
For graphs with cycles, becomes approximate but often works well in practice.

# Summary

- ▶ **General framework** for inference on tree-structured graphs
- ▶ Uses **factor graph** representation for clarity
- ▶ Based on **message passing** between variable and factor nodes
- ▶ Provides **efficient exact marginals** via two-pass algorithm
- ▶ **Foundation** for many important algorithms in machine learning
- ▶ Extends to **approximate inference** for loopy graphs

## The Big Picture

The Sum-Product algorithm transforms the exponentially hard problem of marginalization into a linear-time procedure by exploiting graph structure and local computations, demonstrating the power of graphical models for efficient probabilistic reasoning.

# The Problem: MAP Inference

## Sum-Product vs Max-Product

- ▶ **Sum-Product:** Computes marginal distributions  $P(X_i)$
- ▶ **Max-Product:** Finds most probable configuration

## MAP Inference Problem

Find the assignment that maximizes the joint probability:

$$\mathbf{x}^* = \arg \max_{\mathbf{x}} P(\mathbf{x}) = \arg \max_{\mathbf{x}} \prod_c \psi_c(\mathbf{x}_c)$$

where  $\mathbf{x} = (x_1, x_2, \dots, x_N)$  is a complete assignment.

## Key Idea: Sum → Max

The algorithm is identical to Sum-Product except for one crucial change:

### Sum-Product Message (Factor → Variable)

$$\mu_{f \rightarrow X}(X) = \sum_{\mathbf{Y}} \left( f(\mathbf{Y}, X) \cdot \prod_{Y \in \text{ne}(f) \setminus \{X\}} \mu_{Y \rightarrow f}(Y) \right)$$

### Max-Product Message (Factor → Variable)

$$\mu_{f \rightarrow X}(X) = \max_{\mathbf{Y}} \left( f(\mathbf{Y}, X) \cdot \prod_{Y \in \text{ne}(f) \setminus \{X\}} \mu_{Y \rightarrow f}(Y) \right)$$

$\sum \rightarrow \max$

# Complete Algorithm Steps

## Step 1: Message Passing

- ▶ Identical to Sum-Product structure
- ▶ Choose root node
- ▶ Perform upward pass (leaves to root)
- ▶ Perform downward pass (root to leaves)
- ▶ **But use MAX instead of SUM**

## Step 2: Compute Max-Marginals at Root

For root variable  $X_r$ :

$$P^*(X_r) = \prod_{f \in \text{ne}(X_r)} \mu_{f \rightarrow X_r}(X_r)$$

- ▶  $P^*(X_r)$  is the **max-marginal**
- ▶ Proportional to maximum probability achievable for each value of  $X_r$

## Step 3: Backtrack for MAP Assignment I


### Backtracking Process

1. **Start at root:**

$$x_r^* = \arg \max_{x_r} P^*(x_r)$$

2. **Recurse downward:** For each child, choose value consistent with parent's choice that maximizes incoming messages.

3. **Continue** until all variables assigned.



### Example

## Step 3: Backtrack for MAP Assignment II

For chain  $X_1 \rightarrow X_2 \rightarrow X_3 \rightarrow X_4$ :

- ▶ Choose  $x_4^* = \arg \max \mu_{3 \rightarrow 4}(x_4)$
- ▶ Choose  $x_3^* = \arg \max_{x_3} [P(x_4^* | x_3) \cdot \mu_{2 \rightarrow 3}(x_3)]$
- ▶ Continue to  $x_2^*, x_1^*$

# Max-Sum Algorithm: Log-Domain Version I

## Why Use Log-Domain?

- ▶ Avoids numerical underflow
- ▶ Converts multiplication to addition
- ▶ More numerically stable

## Transformations

Original	Log-Domain
Multiplication	Addition
Maximization	Maximization
$\prod$	$\sum$
max	max

## Max-Sum Messages

# Max-Sum Algorithm: Log-Domain Version II

► **Variable to Factor:**

$$\log \mu_{X \rightarrow f}(X) = \sum_{h \in \text{ne}(X) \setminus \{f\}} \log \mu_{h \rightarrow X}(X)$$

► **Factor to Variable:**  $\log \mu_{f \rightarrow X}(X) =$

$$\max_{\mathbf{Y}} \left[ \log f(\mathbf{Y}, X) + \sum_{Y \in \text{ne}(f) \setminus \{X\}} \log \mu_{Y \rightarrow f}(Y) \right]$$

# Comparison: Sum-Product vs Max-Product

Aspect	Sum-Product	Max-Product
<b>Goal</b>	Compute marginals $P(X_i)$	Find MAP assignment $\mathbf{x}^*$
<b>Operation</b>	Sum over other variables	Max over other variables
<b>Output</b>	Probability distributions	Single best configuration
<b>Complexity</b>	$O(K^m)$ per message	$O(K^m)$ per message
<b>Backtrack</b>	Not needed	Required for MAP assignment
<b>Normalization</b>	Required during/after	Optional (only for probabilities)
<b>Applications</b>	Bayesian inference, estimation	Decoding, segmentation, optimal control